# Governed Feature Selection for Credit-Risk Models: Shapley Attribution with Pre-Registered Parsimony under a Decision-Gate Governance Framework

Ravi*

***Abstract*—Machine-learning credit-risk models are increasingly accurate but remain hard to *govern*: the questions of *which* features are genuinely useful, *how parsimonious* a defensible model may be, and *whether* a selection is adequately powered and fair are usually answered by analyst judgement rather than a reproducible, auditable procedure. We present two coupled contributions. The first is a pre-registered feature-selection methodology that separates order-invariant *attribution* (sampled Shapley importance against a shadow-feature null) from a distinct, fixed *selection* rule (the smallest feature set that is both DeLong-indistinguishable from the full model and within a 5% relative-Gini tolerance), followed by always-on isotonic calibration and a regulatory PD floor. The second is a hash-chained decision-gate governance framework (G0–GF) that wraps it with materiality-tiered rigour, binding every statistical bar to a cited source and separating *governance* decisions from *analytics* parameters, which are never re-opened. Across a corpus of eleven public and reference credit datasets run under one fixed configuration, we show that the same rule yields genuine parsimony at small sample sizes yet near-full feature sets at large ones (a large-$N$ sensitivity effect we characterise), that excluding a protected attribute is necessary but not sufficient to remove disparate impact, and that signal strength—not sample size—sets the achievable discrimination ceiling. We add the Riley *et al.* minimum-sample-size calculation alongside events-per-variable as a reported sufficiency diagnostic and give a retrospective across the corpus.**

***Impact Statement*—Credit decisions affect access to capital for individuals and firms, and supervisors now require model-risk controls (e.g., the U.S. interagency guidance superseding SR 11-7, and the EU AI Act's treatment of creditworthiness assessment as high-risk). This work shows how feature selection—an early, consequential step usually left to discretion—can be made *reproducible and auditable* without inventing economic assumptions: thresholds and pricing remain institution inputs, and governance gates decide process, not analytics. The framework makes a model's parsimony, calibration, fairness, and data-sufficiency claims traceable to recorded runs and cited bars, which should lower the cost of independent validation and reduce the risk of overstated or unfair models reaching deployment.**

***Index Terms*—Credit risk, feature selection, Shapley values, model governance, calibration, fairness, disparate impact, sample size, explainable AI.**

## I. Introduction

MACHINE-learning models now dominate credit-risk scoring, yet their *governance* has not kept pace with their accuracy. A modern scorecard is the product of dozens of consequential choices made before any headline metric is reported—which raw columns to admit, how to bin and screen them, which to keep, how small a model is defensible, and whether the data are even adequate to support it. In practice these choices are made by analyst judgement and documented, if at all, after the fact. That is precisely where model risk concentrates, and precisely what supervisors now scrutinise: the U.S. interagency guidance SR 26-2 (superseding SR 11-7) demands effective, independent challenge of model development; the EU AI Act classifies creditworthiness assessment as high-risk; and Basel CRE36 requires a use test and a per-grade calibration backtest. None of these tells a practitioner *how* to make the early modelling steps reproducible and auditable.

We identify three governance gaps in the feature-selection stage. First, **which features are genuinely useful**: greedy, order-dependent screening produces artefactual retention, and Shapley-style attribution—increasingly the default interpretability tool—fairly apportions importance but does not, by itself, select a subset. Second, **how parsimonious a model may defensibly be**: small differences in discrimination are routinely presented as wins or losses with no inferential test, and the "right" number of features is asserted rather than derived. Third, **whether the data are adequate and the model fair**: sample-size justification rarely goes beyond a rule of thumb, and the removal of a protected attribute is often assumed to remove disparate impact. Each gap is an opening for an overstated or unfair model to reach deployment with a clean-looking paper trail.

We address these with two coupled artefacts. The first is a *pre-registered feature-selection methodology* that cleanly separates order-invariant *attribution* from a fixed *selection* rule, tests feature signal against an empirical shadow-feature null, calibrates by construction, and reports both events-per-variable and the Riley minimum sample size as governed sufficiency diagnostics. The second is a *decision-gate governance framework* that wraps the methodology in a materiality-tiered sequence of gates, each closing with an accountable human sign-off written to a tamper-evident hash chain, and—by design—never lets a per-run governance decision re-open a pre-registered analytical parameter. Evaluated as a single fixed configuration across eleven credit datasets, the pair exposes behaviour that discretion would hide: the same selection rule yields genuine parsimony at small $N$ but near-full feature sets at large $N$; excluding a protected attribute does not remove disparate impact; and signal strength, not sample size, bounds achievable discrimination.

**Contributions.**

1) A **pre-registered feature-selection methodology** separating Shapley *attribution* from a fixed *parsimony* selection rule (§III).
2) A **hash-chained, materiality-tiered decision-gate governance framework** binding every bar to a cited source and separating governance from analytics (§IV).
3) An empirical study over eleven datasets under one fixed configuration, including the large-$N$ parsimony effect, the "necessary-but-not-sufficient" fairness finding, and a Riley-*et al.* sample-size retrospective (§VI).

## II. Related Work

**Feature importance and selection.** Classical selection splits into filter methods (e.g. Information Value, the credit-scoring standard [1]), wrapper methods, and embedded penalisation (LASSO, elastic net). Shapley-based attribution has become the dominant interpretability tool: SHAP [2] gives local, per-prediction attributions, while global measures such as SAGE [3] and the Shapley Population Variable Importance Measure (SPVIM) [4] attribute model-level

*The author is with the Indian Institute of Foreign Trade, New Delhi, India (e-mail: ravi_phdm23@iift.edu).

Manuscript prepared October 18, 2026. This is a working draft for internal review, not a submitted version.

predictiveness. A recurring confusion—which our methodology is built to avoid—is treating these attributions as a *selection* criterion: a Shapley value fairly *apportions* predictiveness but does not identify a minimal sufficient subset, so an importance ranking must be paired with an explicit selection rule. Boruta's shadow-feature device [5] and stability selection [6] address the related problem of distinguishing real signal from noise; our methodology adopts the shadow-null idea as its retention test.

**Parsimony and equivalence.** Preferring the smallest model statistically consistent with the best one has a long heritage—the one-standard-error rule of CART [7] and its treatment in the statistical-learning literature [8]. Our methodology operationalises this with a paired-AUC equivalence test [9] plus a relative-Gini tolerance, rather than an information criterion, because discrimination is the quantity credit practice cares about [10], [1].

**Model-risk governance.** Supervisory expectations for model risk—the U.S. interagency guidance SR 26-2 [11] (superseding SR 11-7), Basel CRE36 [12] (use test, per-grade calibration backtest), the EU AI Act's [13] high-risk regime for creditworthiness assessment, and the EBA's Margin-of-Conservatism guidance—specify *what* must be validated, documented, and independently challenged, but not a concrete, reproducible procedure that couples feature selection to signed, tamper-evident approval. Our governance framework fills that gap by turning each analytical stage into a governed gate with a hash-chained decision log. Documentation artefacts such as model cards [14] and datasheets for datasets [15] standardise *what* to report about a model or data source, but they are descriptive records rather than a gated, append-only decision process; our framework is complementary, producing such records as a by-product of signed gate decisions.

**Fairness and disparate impact.** Anti-discrimination practice distinguishes disparate *treatment* from disparate *impact*, the latter commonly screened with the four-fifths rule (EEOC [16]). A substantial fairness-in-ML literature formalises group-fairness criteria [17]; a consistent empirical lesson, which our runs reproduce, is that removing a protected attribute is necessary but not sufficient, because correlated features reconstruct it.

**Sample-size adequacy.** Credit-scoring studies rarely report a sample-size justification beyond the events-per-variable rule of thumb [18]. Recent work shows EPV is too coarse [19] and derives signal-aware minimum sample sizes for prediction models [20], [21], with a C-statistic-to-$R^2$ route for forward use [22]. Our methodology reports both EPV and the Riley minimum-$N$ as governed diagnostics. In sum, each ingredient exists in isolation; our contribution is their coupling into one pre-registered, reproducible, and auditable pipeline, evaluated here as a fixed-config corpus.

## III. The Feature-Selection Methodology

The proposed methodology deliberately separates two questions that are routinely conflated in applied feature selection (the full pipeline and its governance gates are shown in Fig. 1):

- **Attribution** — "How much predictive value is attributable to this feature, *order-invariantly*?" Answered by a Shapley importance measure against an empirical null.
- **Selection** — "What is the *smallest* feature set whose performance is statistically indistinguishable from the full set?" Answered by a separate, pre-registered stopping rule.

The separation is principled, not stylistic. Shapley values are a *fair-attribution* solution concept (efficiency, symmetry, dummy, additivity); they do not identify an optimal subset. A redundant-but-informative feature still receives a positive share, and a feature whose value concentrates in rare coalitions can average below any bar despite being needed. Attribution therefore *ranks*; a separately justified rule *selects*. Greedy, category-sequential screening—the common alternative—produces order-dependent, artefactual retention; an order-invariant importance estimate plus an explicit selection rule avoids it.

**Algorithm 1** Governed feature selection (one dataset, fixed configuration)

**Require:** data $(X, y)$; learner; value function $v$; shadows $m \geq 10$; tolerance $\tau$=0.05; level $\alpha$=0.05; confirmation threshold $c$=0.80; seed
screen features by $\mathrm{IV} \geq 0.02$ with monotonic WOE binning ▷ Stage 1
group columns at $|\rho| \geq 0.7$ into coalitional players $P$ ▷ Stage 2
add $m$ shadow players: row-permuted copies of real columns
**if** $|P| + m \leq 15$ **then**
    compute exact Shapley values $\phi$ for all players
**else**
    estimate $\phi$ by sampled permutations to a target MC s.e.
**end if**
$b \leftarrow$ 95th percentile of the shadow players' $\phi$ ▷ empirical null
$C \leftarrow \{\, p : \phi_p > b$ and bootstrap confirm. freq $\geq c \,\}$
rank $C$ by $\phi$ descending ▷ Stage 3
**for** $k = 1, \ldots, |C|$ **do**
    fit top-$k$ on out-of-fold folds; record $\mathrm{Gini}_k$ and DeLong $p_k$ vs full
**end for**
$k^\star \leftarrow \min\{\, k : p_k \geq \alpha$ and $(\mathrm{Gini}_{\text{full}} - \mathrm{Gini}_k)/\mathrm{Gini}_{\text{full}} \leq \tau \,\}$
calibrate the top-$k^\star$ model (isotonic); apply PD floor/cap ▷ Stage 4
**Ensure:** selected set (top-$k^\star$ of $C$), calibrated model, additive reason codes

**Inference target.** We estimate population feature importance *with respect to a specified learner and value function*: how much cross-validated predictiveness is recoverable from a feature set, for a given learner class, on a given population—*not* how much a feature "matters to credit risk" in the abstract. Formally, with value function $v(S)$ equal to the cross-validated predictiveness of the columns in $S$, the importance of feature $j$ is the Shapley value

$$\phi_j = \sum_{S \subseteq N \setminus \{j\}} \frac{|S|!\,(|N| - |S| - 1)!}{|N|!} \big[ v(S \cup \{j\}) - v(S) \big], \quad (1)$$

which is the Shapley Population Variable Importance Measure (SPVIM) of Williamson and Feng [4]. This is *not* SAGE [3], which marginalises absent features out of a single fitted model; here each coalition $S$ is *refit*.

The pipeline has five stages (0–4), summarised in Algorithm 1; the governance framework that signs off on their outputs is described in §IV.

### A. Stage 0: Eligibility and data governance

Stage 0 runs *before any statistics* and drops no column; it attaches a per-column classification that travels with every downstream result. Each column is tagged `permissible`, `prohibited` (e.g. sex, age, marital status), `proxy_risk` (e.g. geography, education, household composition), or `unclear`. On these public research benchmarks, protected and proxy attributes are *retained with mandatory disclosure* so results remain comparable with the large published literature on the same datasets; no model here is intended for live use. Stage 0 also records point-in-time availability (flagging `near_target` and `external_score` columns that make a model a behavioural or score-recovery exercise rather than an application scorecard), audits sentinel codes and shared missingness for leakage, notes the accept-only sampling that truncates feature distributions [23], and declares *mandatory* columns and a missingness policy that are enforced only at scoring (§III-F).

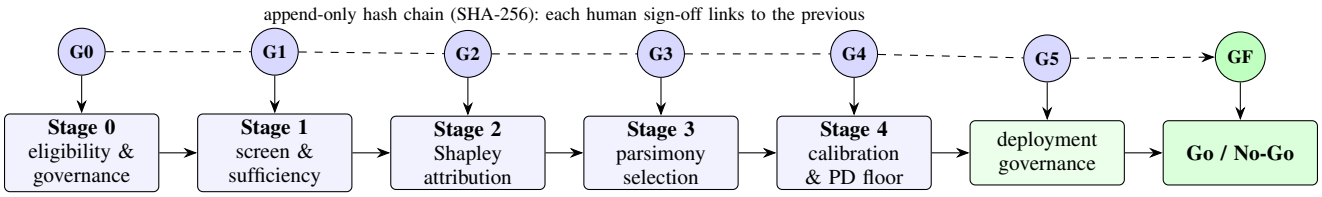



Fig. 1. The governed pipeline. The feature-selection methodology (blue, Stages 0–4) runs under a sequence of governance gates (G0–G5), each closed by an accountable human sign-off that is written to an append-only hash chain; the committee gate GF rules the final Go/No-Go with all gate evidence visible. Gates decide scope, eligibility, fairness treatment, cut-offs and disposition—never the pre-registered analytical parameters (P8).

### B. Stage 1: Univariate screening

Information Value (IV) is the single a-priori screen, at the conventional threshold IV $\geq 0.02$ [1]; Cramér's $V$, $\chi^2$ and $p$-values are reported descriptively and decide nothing. Binning is monotonic Weight-of-Evidence (WOE) with a minimum bin size of 5% of rows,

$$\mathrm{WOE}_i = \ln \frac{g_i/G}{b_i/B}, \qquad \mathrm{IV} = \sum_i \Big(\tfrac{g_i}{G} - \tfrac{b_i}{B}\Big)\mathrm{WOE}_i, \qquad (2)$$

where $g_i, b_i$ are goods/bads in bin $i$. Monotonicity is enforced only where it is supported: a genuinely U-shaped feature whose IV recovers $\geq$ 20% when the constraint is freed is tagged `shape_nonmonotonic` and routed to a single-column model check rather than being judged on its monotonic IV. A reliability gate ($\chi^2$ minimum expected cell $\geq 5$) and a bin-count sensitivity check (WOE vs 5/10/20 quantile bins) escalate borderline columns to the same model check, whose verdict is a Gini confidence interval excluding zero.

### C. Stage 2: Order-invariant attribution

The default value function is repeated cross-validated AUC and the default learner is a penalised logistic model on monotonic-WOE inputs (`woe_logistic`)—the credit-scoring standard form: deterministic and hyperparameter-insensitive, which removes refit noise from the Shapley estimate. XGBoost [24] is run as a robustness learner; material disagreement in the two rankings (reported via Kendall's $\tau$) is itself a finding about interaction dependence. Class reweighting is *off* by default: on GMSC it inflated the mean predicted probability to $0.286$ against a true base rate of $0.067$, and the resulting miscalibration is avoided rather than repaired.

Two devices make the attribution both tractable and honest. First, correlated columns are not pruned (which would reintroduce a greedy univariate criterion); instead, hierarchical clustering on $1 - |\rho_{\text{Spearman}}|$ cut at $|\rho| = 0.7$ makes *each cluster a single Shapley player*, preserving the efficiency property over the full candidate set. Second—the central methodological fix—significance testing of Shapley values is replaced by a **shadow-feature null** [5]: at least ten shadow players (row-permuted copies of real columns) are added, the identical Shapley computation is run over the augmented set, and a real player is `importance_confirmed` only if its value exceeds the $95^{\text{th}}$ percentile of the shadow distribution *and* its bootstrap confirmation frequency is $\geq 0.80$. The bar is a quantile of an empirical null, so no multiple-testing correction is needed and fold-dependence does not bias it. Benjamini–Hochberg [25] $p$-values and Cohen's $d$ [26] are retained only as reported diagnostics—usefully self-refuting, since on German Credit a permuted-noise shadow column received $p = 0.0000$ with a *negative* Shapley value, which a significance rule would have confirmed and the shadow bar correctly rejected. Enumeration is exact when the player count (including shadows) is $\leq 15$ and sampled otherwise, with the Monte Carlo standard error reported per player.

### D. Stage 3: Pre-registered parsimony

The eligible set is the confirmed players, ranked by $\phi_j$. For $k = 1, \ldots, K$ the model on the top-$k$ players is scored on *out-of-fold* predictions (so every $k$ is evaluated on the same rows and DeLong's paired test is valid [9]), and the full parsimony curve is a required output. The stopping rule is fixed in advance: choose the smallest $k$ that is both statistically indistinguishable from the full-candidate model by DeLong's test at $\alpha = 0.05$ *and* within a relative-Gini tolerance,

$$k^\star = \min\Big\{k : \; p_{\text{DeLong}}(k) \geq 0.05 \;\wedge\; \tfrac{\text{Gini}_{\text{full}} - \text{Gini}_k}{\text{Gini}_{\text{full}}} \leq 0.05\Big\}. \qquad (3)$$

If no $k$ qualifies, that is reported as the result ("no parsimonious subset meets the tolerance")—it is never replaced by the best available $k$. Because the two conditions are combined by *and*, the rule behaves very differently across sample sizes (§VI): at small $N$ both clear together and yield genuine parsimony, whereas at large $N$ DeLong detects the loss from dropping even a weak feature and the full set is returned.

### E. Stage 4: Calibration, reason codes, and sufficiency diagnostics

The selected model is calibrated by cross-validated isotonic regression [27], [28] (always on), and calibrated probabilities are used whenever a PD is quoted to a person. Because the learner is WOE-logistic, the score is an exact additive decomposition $\text{logit} = \beta_0 + \sum_j \beta_j \,\mathrm{WOE}_j(x)$, giving faithful per-feature *reason codes* with no post-hoc approximation. A regulatory PD floor of $0.03\%$ (Basel CRE32.3) and a $99.9\%$ cap are applied at scoring, as are the Stage-0 mandatory-field gate and missingness policy. Finally, data sufficiency is reported—never used to auto-reject—through two diagnostics side by side: events-per-variable [18], [19] and the Riley *et al.* minimum sample size [21], [20], the latter derived from the anticipated $C$-statistic and prevalence via the Cox–Snell $R^2$ approximation of [22]. A shortfall is a disclosed caveat the governance layer weighs, not a gate failure—consistent with the principle that selection parameters are fixed and only governance decides disposition.

## IV. The Governance Framework

The methodology of §III produces a defensible *model*; it does not, by itself, produce a defensible *decision to use that model*. A supervisor or internal audit function asks questions the analytics cannot answer: who decided the data were adequate, and could they have decided otherwise? Who chose this feature set, on what recorded rationale? Who is accountable if the model is wrong, and were they independent of the builder? Where is the evidence that each regulatory expectation was *met*, not merely claimed? The governance framework answers these by turning each methodology stage into a *governed decision gate*: the machine emits a recommendation and a full evidence pack, a human with the right role and sufficient independence chooses from a closed set of genuinely viable options, and the choice is written to a tamper-evident log. This mirrors modern

TABLE I
THE DECISION GATES AND THE METHODOLOGY STAGE EACH GOVERNS.

| Gate | Stage | Core question decided |
|---|---|---|
| G0 | Stage 0 | Is the dataset fit to build on? (coverage, label integrity, scope) |
| G1 | Stage 1 | Which features survive screening; is $N$ adequate (EPV + Riley, reported)? |
| G2 | Stage 2 | Is the attribution stable enough to select on? |
| G3 | Stage 3 | Which feature set is adopted—parsimonious vs full, and the trade-off? |
| G4 | Stage 4 | Calibrated, discriminating, and fair enough to approve? (per-grade backtest) |
| G5 | Scoring | Deployment controls + cutoff policy + acceptance demonstration |
| GF | whole model | GO / GO-WITH-CONDITIONS / NO-GO (committee) |

model-risk expectations—the U.S. interagency guidance SR 26-2 (superseding SR 11-7), the EU AI Act's treatment of creditworthiness assessment as high-risk, and Basel CRE36—without importing their numeric bars by assumption: every threshold used is tied to a cited source.

### A. Design principles

Eight principles govern the framework; three are load-bearing. **(P1) Machine recommends, human decides, record proves**: every gate is recommendation + evidence pack + closed options + accountable sign-off + log entry. **(P2) Segregation of duties**: the Develop, Validate, and Approve functions are never held by one person for one model. **(P8) Gates decide governance, not analytics**: a gate's option set may *never* re-open a pre-registered analytical parameter (estimator mode, selection tolerances, shadow-bar quantile, confirmation frequency, `random_state`, CV folds). Those are fixed in the methodology, identical across every dataset, and chosen once; gates decide scope, eligibility, fairness treatment, cutoffs, and Go/No-Go. P8 is what keeps the method reportable: letting a per-run gate switch the estimator would be the dataset-dependent, path-changing behaviour a research methodology must avoid. The single failure mode the framework is explicitly designed against is *rubber-stamping*—hence the anti-strawman rule below and the independence requirements in §IV-E.

### B. The decision gates G0–GF

The methodology's stages become gates G0–G5, with a final committee gate GF (Table I). G0–G5 are *stage gates* (a single accountable role signs each); GF is the *committee gate*, ruled by the Model Approval Committee. Crucially, GF is last and is decided only after G4's per-grade calibration backtest and G5's cutoff policy and acceptance demonstration are complete and visible—one does not approve deployment and then check whether the PDs are honest and the decisions sensible.

A model is assigned a *materiality tier* (HIGH / MEDIUM / LOW) at G0, after SR 26-2's exposure-$\times$-purpose logic; the tier sets *who* signs each gate and *how high the pass-bar sits*, but never changes the analytics—every metric is measured identically across datasets. Lowering a tier is itself a recorded G0 decision.

### C. Anatomy of a gate

Each gate produces one Gate Decision Record with five parts: (i) the machine recommendation and its numeric basis; (ii) the full evidence pack (the actual numbers, traceable to a run, not a summary); (iii) a closed set of 2–4 options, each carrying its own data, rationale, benefits, risks, and an explicit *residual-risk acceptance*; (iv) a human sign-off (role, identity, exactly one option, free-text rationale, timestamp); and (v) an append-only log entry. The *anti-strawman rule* forbids padding: if the machine cannot construct a genuine case for an alternative, it presents fewer options and flags "single viable option; escalate," routing the decision upward rather than manufacturing fake choices.

TABLE II
EXAMPLE GATE DECISION RECORD—THE CLOSED, VIABLE OPTION SET AT G3 (SELECTION) FOR CREDITRISK, EACH WITH THE RESIDUAL RISK THE SIGNER ACCEPTS. THE MACHINE MUST DEFEND EVERY OPTION; UNVIABLE "STRAWMAN" OPTIONS ARE NOT PRESENTED.

| Option | Key data | Residual risk accepted |
|---|---|---|
| **A** (chosen): ratify $k$=6 benchmark | Gini 0.76; DeLong $p$=0.88; rel. loss $-0.01\%$ | benchmark carries circular features; not origination-valid |
| **B**: origination-only 4-feature model | Gini 0.61 ($-19\%$); independent of incumbent | lower discrimination for application-time use |
| **C**: NO-GO | — | delay; a solid independent signal exists, so unwarranted |

Tamper-evidence is a hash chain over the decision records. With a genesis token $h_0 = \mathsf{GEN}\text{-}\langle\text{TAG}\rangle\mathsf{0}$, each gate's entry hash is

$$h_i = \text{SHA-256}\big(h_{i-1} \,\|\, c_i\big)[0{:}8], \tag{4}$$

where $c_i$ is the gate's decision content and $\|$ is concatenation. Any later edit to a sealed decision breaks every subsequent link, so a reversal must be a *new* recorded decision, never a silent edit—the audit trail is append-only. Table II shows a representative decision record: the closed, viable option set at the selection gate for CREDITRISK, each option carrying the residual risk the signer explicitly accepts.

### D. Roles, segregation of duties, and the committee

The framework defines seven roles; the hard rule (P2) is that Develop, Validate, and Approve are never the same person for one model, mirroring SR 26-2's conflict-of-interest provisions. Validators sign G1–G4 and can *return* a gate; **G4 requires dual sign-off** (Risk Officer *and* Validator), because calibration and fairness are where model error becomes capital error and conduct risk. The Model Approval Committee—chaired by the CRO or an independent delegate, with a quorum that is void without an independent validator present—rules GF. Any member who developed the model recuses and is minuted; the developer presents but does not vote. Every MAC decision is itself a hash-chained event.

### E. The Go/No-Go engine

At GF a deterministic engine synthesises six inputs—dataset adequacy, the signed gate decisions, residual risk vs appetite, the regulatory checklist, audit expectations, and the new-model acceptance bar—into a recommendation the committee rules on: *all gates signed and every regulatory criterion satisfied and residual risk within appetite* $\Rightarrow$ **GO**; *remediable partials with conditions* $\Rightarrow$ **GO-WITH-CONDITIONS** (the engine lists the conditions); *any gap on a must-have criterion, or an unsigned/returned gate, or residual risk outside appetite* $\Rightarrow$ **NO-GO** (the engine names the blocking item and the gate to return to). The engine never upgrades a NO-GO; the committee may downgrade on judgement and must minute why, but cannot clear a hard gap except by resolving it as a new, re-synthesised gate decision.

## V. Experimental Setup

**Datasets.** The corpus is eleven binary-outcome credit datasets (Table III): widely used public benchmarks—including the UCI German [29], [30] and Taiwan [31] credit datasets, the Home Equity (HMEQ) [32] and Give-Me-Some-Credit [33] sets, the FICO HELOC panel [34], a LendingClub extract [35], and the PAKDD-2009 Brazilian retail-credit set [36]—alongside additional reference portfolios. They span $N \in [1{,}000, 149{,}391]$, prevalence $[6\%, 52\%]$, and candidate pools of 9–87 features. Per Stage 0, protected and proxy attributes are retained with disclosure so results remain comparable with the published literature on these datasets; no model here is intended for live use. Full provenance and target definitions are recorded per dataset.

**Fixed configuration (identical for every dataset, by P8).** Primary learner `woe_logistic` (XGBoost as a robustness learner); value function cross-validated AUC; correlated columns grouped at $|\rho| = 0.7$; $\geq 10$ shadow players with a $95^{\text{th}}$-percentile bar and confirmation frequency $\geq 0.80$; exact Shapley when players (including shadows) $\leq 15$, sampled otherwise; selection rule (3); isotonic calibration; PD floor $0.03\%$ and cap $99.9\%$. The random seed is 42 throughout. No gate re-opens any of these.

**Evaluation protocol.** Splits are stratified $80/20$ taken *before* any resampling, with the test set held at its natural base rate; the parsimony curve is scored on out-of-fold predictions from stratified 5-fold cross-validation so DeLong's paired test is valid. Discrimination is reported as Gini (primary) with AUC, KS, partial AUC, Brier and ECE [37]; paired model comparisons use DeLong [9]; calibration uses Hosmer–Lemeshow [38] and a per-grade realised-vs-expected backtest; fairness uses the adverse-impact ratio against the four-fifths rule; data sufficiency uses EPV [18] and the Riley minimum-$N$ [21].

**Governance and reproducibility.** Each dataset is run end-to-end through gates G0–GF with a human sign-off recorded at every gate and appended to the hash chain (4); the materiality tier is set at G0. The pipeline is deterministic: independent re-runs of two datasets reproduced their committed results exactly—German (Gini 0.53, GO-WITH-CONDITIONS) including its hash chain, and CREDITRISK (selection Gini 0.76), the latter re-run on a *cleaned* copy of the file (its corrupt label column removed) yet returning the identical screen, attribution, and $k$=6 selection, confirming that the result depends only on the modelling inputs. Validation is i.i.d. cross-validation only—no dataset has a reliable time field, so no claim of out-of-time or production readiness is made—and all accept/decline thresholds and pricing are treated as exogenous institution inputs, never derived here.

## VI. Results

We run eleven public and reference credit datasets through the full governed exercise (G0–GF) under *one fixed configuration* (seed 42; selection = smallest $k$ with DeLong $p \geq 0.05$ and $\leq 5\%$ relative-Gini loss; sampled Shapley when players $> 15$). The corpus spans $N$ from 1,000 to 149,391, prevalence from $6\%$ to $52\%$, and candidate pools from 9 to 87 features. Table III and Fig. 2 summarise it; nine runs seal GO-WITH-CONDITIONS and two (MSME, PAK) seal NO-GO. Every figure traces to a recorded run.

### A. A worked run: German credit, G0 to GF

To make the pipeline concrete, we trace one dataset end-to-end. On German credit ($N$=1,000, 30% default rate), **G0** excludes the three protected attributes (age, sex/marital status, foreign-worker status) and sets a MEDIUM materiality tier, leaving a 17-feature pool. **G1** accepts the univariate screen—12 of 17 features pass, and both sufficiency diagnostics are met (EPV ≈17.6; Riley minimum 323 vs 1,000 available)—with `checking_status` dominant (IV 0.67).

TABLE III
THE GOVERNED CORPUS. POOL = FEATURES ENTERING THE PIPELINE AFTER STAGE-0; SEL = SELECTED; GINI/AUC ARE OUT-OF-FOLD AT SELECTION. VERDICT IS THE SEALED GF DECISION.

| Dataset | $N$ | base | pool→sel | Gini | AUC | Verdict |
|---|---|---|---|---|---|---|
| HMEQ | 5,960 | .200 | 12→11 | 0.81 | 0.91 | GO-cond |
| CREDITRISK | 32,581 | .218 | 9→6 | 0.76 | 0.88 | GO-cond |
| GMSC | 149,391 | .067 | 9→8 | 0.70 | 0.85 | GO-cond |
| HELOC | 10,459 | .522 | 23→17 | 0.59 | 0.79 | GO-cond |
| TAIWAN | 30,000 | .221 | 20→14 | 0.54 | 0.77 | GO-cond |
| GERMAN | 1,000 | .300 | 17→5 | 0.53 | 0.77 | GO-cond |
| MSME | 1,707 | .060 | 87→15 | 0.43 | 0.72 | **NO-GO** |
| LC | 39,717 | .142 | 21→9 | 0.36 | 0.68 | GO-cond |
| LOAN | 9,578 | .160 | 13→7 | 0.34 | 0.67 | GO-cond |
| TH02 | 1,225 | .264 | 14→6 | 0.27 | 0.64 | GO-cond |
| PAK | 50,000 | .261 | 39→6 | 0.15 | 0.57 | **NO-GO** |

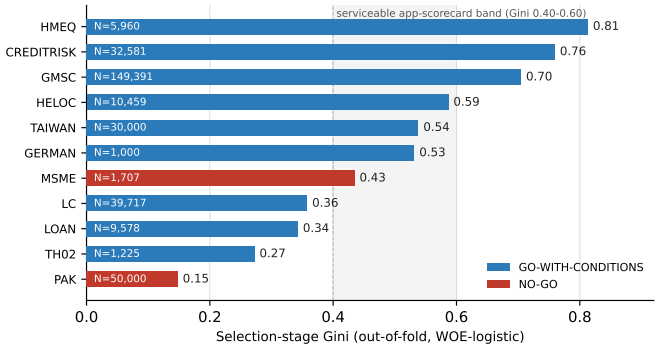



Fig. 2. Selection-stage Gini across the governed corpus, coloured by sealed GF verdict. Discrimination spans an order of magnitude; the two NO-GO runs are not simply the lowest-Gini—MSME is rejected on sufficiency, not discrimination.

**G2** confirms 10 of 12 players above the shadow-feature null; the U-shaped `credit_amount`, which a monotonic IV screen underrates, is vindicated at the model check. **G3** selects $k$=5 features (`checking_status`, `credit_history`, `months_duration`, `savings_status`, `purpose`), Gini 0.53—the smallest $k$ at which the equivalence test ($p$=0.06) and the 5% tolerance (4.1%) clear *together*. **G4** finds calibration sound (Hosmer–Lemeshow $p$=0.46, ECE 0.04) and fairness mixed: the sex adverse-impact ratio is clean (0.99) while the age ratio fails four-fifths (0.57) *despite age being excluded*—a true-risk-driven disparity reconstructed through correlated features. **G5** fixes the scoring controls (mandatory `checking_status`; exact reason codes; PD floor), and **GF** rules GO-WITH-CONDITIONS. Every decision is signed and appended to the hash chain, which seals at a final digest; an independent re-run reproduced every gate and the chain exactly (§V).

### B. Parsimony varies with sample size

The single fixed stopping rule (3) produces strikingly different parsimony depending on $N$, because DeLong's power to detect the loss from dropping a weak feature grows with sample size. At **small** $N$ the rule genuinely parsimonises: on German ($N$=1,000) it selects 5 of 17 features, with both conditions clearing together (DeLong $p$=0.061 $\geq 0.05$ and relative-Gini loss $4.10\% \leq 5\%$). At **large** $N$ the same rule returns a near-full set: on LC ($N$≈40,000) it keeps 8 of 9 players (DeLong $p$=1.0, loss 0.0%), and on HMEQ ($N$≈6,000) DeLong rejects *every* proper subset ($k$=9 vs full, $p = 6.1 \times 10^{-4}$) even though a 5%-loss criterion alone would admit a six-feature model—so the *and*-rule returns all eleven. This is the methodology declining to overstate parsimony when the data can resolve the difference. We

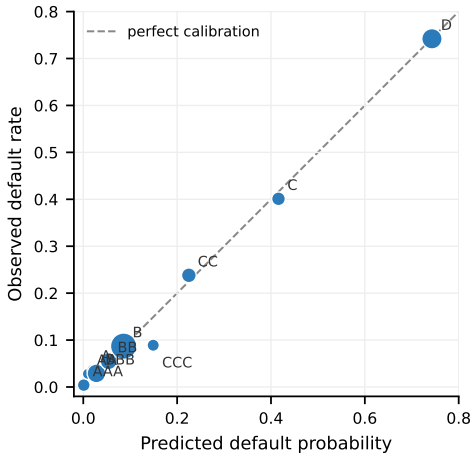



Fig. 3. Per-grade calibration reliability on the CREDITRISK held-out set (rating grades AAA–D; marker area $\propto$ grade population). Points lie close to the diagonal across the full risk range; Hosmer–Lemeshow is not rejected ($p$=0.44, ECE 0.011). The one mild outlier is the small CCC grade, which over-predicts (a conservative direction).

call this a *large-sample significance effect*: the same rule, opposite behaviour at the two ends of the sample-size range, both honest. A corollary visible in Table III is that *signal, not volume, sets the ceiling*—German reaches Gini 0.53 on 1,000 rows while LC reaches 0.36 on 40,000.

### C. Fairness: excluding a protected attribute is necessary but not sufficient

On two independent datasets, removing the protected attribute did not remove disparate impact, because correlated features reconstruct it. On **LC**, geography (`addr_state`) was dropped at a cost of only 0.007 Gini, yet a residual regional approve-rate disparity persisted (worst adverse-impact ratio 0.62, below the four-fifths bar). On **German**, `age` was excluded at G0, yet the age-band approve-rate AIR was 0.57 (under-25 default rate 48.7% vs 13.3% for 60+); because the model is accurate *within* each band, this is disparate *impact driven by true risk*, not proxy bias. Both runs therefore record the same governed finding—proxy removal is necessary but not sufficient, and a fairness-constrained replication or reject-inference study would be required before any live use—surfaced as a binding GF condition rather than discovered post-deployment.

### D. Calibration

With class reweighting off and isotonic calibration always on, calibrated probabilities are tight across the corpus: expected calibration error is 0.020 on HMEQ, 0.012 on LC, and 0.040 on German, and Hosmer–Lemeshow is *not* rejected even at the largest $N$ (GMSC, LC)—the feared large-$N$ hypersensitivity did not materialise for calibration, only for parsimony. Because the learner is WOE-logistic, each PD carries an exact additive reason-code decomposition. Fig. 3 illustrates the per-grade reliability on CREDITRISK, whose model spans the full AAA–D range: observed default rates track the predicted probabilities closely across nearly two orders of magnitude.

### E. Sample-size adequacy (EPV and Riley)

Table IV reports the two sufficiency diagnostics. Ten of eleven runs meet the Riley minimum; the lone shortfall is **MSME** (requires 3,654, has 1,707), which is *also* the only run failing EPV ($6.8 < 10$)

TABLE IV
RILEY MINIMUM-$N$ RETROSPECTIVE (REPORTED DIAGNOSTIC).
$N_{\text{REQ}} = \max(B1, B3, B4)$.

| Dataset | $N$ | $N_{\text{req}}$ | EPV | Meets? |
|---|---|---|---|---|
| HMEQ | 5,960 | 278 | 108 | yes |
| CREDITRISK | 32,581 | 263 | 1185 | yes |
| GMSC | 149,391 | 570 | 1251 | yes |
| HELOC | 10,459 | 522 | 321 | yes |
| TAIWAN | 30,000 | 737 | 474 | yes |
| GERMAN | 1,000 | 323 | 60 | yes |
| MSME | 1,707 | 3,654 | 6.8 | **SHORT** |
| LC | 39,717 | 1,581 | 625 | yes |
| LOAN | 9,578 | 1,220 | 219 | yes |
| TH02 | 1,225 | 1,170 | 54 | yes (+55) |
| PAK | 50,000 | 4,041 | 2173 | yes |

and the only HIGH-tier NO-GO—the two diagnostics are concordant, and the single failure is the case already rejected. Riley adds what EPV cannot see: **PAK** has an enormous EPV (2,173) yet needs 4,041 rows, because its weak signal (AUC 0.57 $\Rightarrow R^2_{cs}$=0.013) makes the shrinkage criterion demand volume; and **TH02** passes by only +55 rows, a thinness EPV (54) does not convey.

## VII. DISCUSSION

**Governance without contaminating the analytics.** The central design choice—that gates decide scope, eligibility, fairness treatment and disposition but never the analytical method (P8)—is what makes the corpus comparable at all. Because the estimator, tolerances, shadow bar and seed are fixed across every dataset, a difference in outcome is a property of the data, not of a per-run tuning decision. The governance layer buys traceability and accountability; it does not buy, and must not buy, the freedom to change the method until the answer improves.

**What the findings mean for practice.** This large-sample effect is not a defect of the rule but a faithful consequence of it: at large $N$ a paired-AUC test can resolve the loss from dropping even a weak feature, so an honest equivalence criterion returns a fuller model. Practitioners should therefore expect genuine parsimony mainly at modest sample sizes and treat a large-$N$ "parsimonious" claim with suspicion. The fairness result sharpens a familiar warning: because correlated features reconstruct an excluded attribute, proxy removal is a floor, not a remedy, and the responsible posture is to *measure and disclose* residual disparate impact (here as an indicative adverse-impact ratio, with the accept/decline cutoff left exogenous) and bind a fairness-constrained follow-up, rather than to claim fairness by omission.

**Limitations.** This study validates a *procedure*, not a performance claim: we do not yet assert that the methodology selects better subsets than established alternatives—that requires the pre-specified benchmark against LASSO, elastic net, RFE, Boruta and TreeSHAP, which is implemented but not run across the corpus. Validation is i.i.d. cross-validation; no dataset has a reliable time field, so no out-of-time or production-readiness claim is made. All datasets are accept-only samples, so feature effects are subject to range restriction [23]. The governance roles are exercised by a single researcher in this methodological setting; the framework specifies the *structure* of independent challenge rather than demonstrating a multi-party deployment. Finally, the Riley diagnostic uses the selected-feature count as its parameter count and a closed-form approximation that can under-state the requirement at very high $C$; neither affects any verdict here, but both are disclosed.

## VIII. Future Work

Several directions follow directly from the limitations above. *Comparative evaluation.* The pre-specified benchmark against LASSO, elastic net, recursive feature elimination, Boruta and TreeSHAP—implemented but not yet run across the corpus—will test whether the methodology selects *better* subsets than established alternatives, upgrading the present "defensible" claim to a comparative one. *Temporal validation.* Applying the pipeline to a dataset with a reliable time index would enable out-of-time validation and a survival / competing-risks extension, replacing the current i.i.d. cross-validation scope and supporting a production-readiness claim. *Reject inference.* Because all datasets are accept-only, incorporating reject inference would correct the range restriction that attenuates feature effects, and would sharpen the fairness analysis in particular. *Fairness-constrained selection.* Since excluding a protected attribute does not remove disparate impact, a natural extension searches for a less-discriminatory alternative—through fairness-constrained selection or a penalised objective—and reports the discrimination cost of each constraint, making the fairness trade-off an explicit, governed decision. *Multi-party governance.* Exercising the role separation with genuinely independent signers, integrated with a live model-risk workflow, would let us measure the governance overhead and the framework's effect on validation turnaround. *Sample-size refinement.* For high-discrimination settings where the closed-form Riley bound can under-state the requirement, a simulation-based minimum sample size and a candidate-pool parameter count would tighten the reported sufficiency diagnostic.

## IX. Conclusion

We have presented a pre-registered feature-selection methodology that separates Shapley attribution from an explicit parsimony rule and reports its own calibration, fairness and sample-size adequacy, wrapped in a hash-chained, materiality-tiered governance framework that records who decided what, on what evidence, and under which cited bar. Run as one fixed configuration across eleven credit datasets, the pair surfaces findings that analyst discretion would obscure—sample-size-dependent parsimony, disparate impact that survives protected-attribute removal, and the primacy of signal over volume—while inventing no economic assumptions: thresholds and pricing remain institution inputs. We hope the reproducible, auditable discipline demonstrated here lowers the cost of independent model validation and raises the bar for what counts as a defensible credit model.

## Appendix: Reproducibility

The configuration is fixed across every dataset and chosen once (Table V); no governance gate re-opens any of it (P8). Tamper-evidence is an append-only hash chain over the signed gate records: with genesis token $h_0$, each entry is $h_i = \text{SHA-256}(h_{i-1} \,\|\, c_i)[0{:}8]$ for decision content $c_i$ (Eq. 4). The sealed chain for the German run passes from genesis `GEN-GERMAN0` through `b1d70c29`, `5139746d`, `2d292965`, `49559806`, `403d8a81`, `884c20a1` to the final digest `a6adf8b1`; any later edit to a sealed record changes its digest and breaks every subsequent link, and an independent re-run reproduced the chain exactly.

TABLE V
The fixed configuration (identical for every dataset; seed 42).

| Parameter | Value |
|---|---|
| Primary / robustness learner | WOE-logistic / XGBoost |
| Value function | cross-validated AUC |
| Correlation grouping | single player at $\lvert\rho\rvert \geq 0.7$ |
| Shadow players | $\geq 10$; bar $=$ 95th percentile |
| Confirmation frequency | $\geq 0.80$ |
| Exact vs sampled Shapley | exact if players $\leq 15$, else sampled |
| Selection rule | smallest $k$: DeLong $p \geq 0.05$ *and* rel. Gini loss $\leq 5\%$ |
| Calibration | isotonic, 5-fold |
| PD floor / cap | $0.03\%$ / $99.9\%$ |
| Split | stratified $80/20$ pre-resampling; 5-fold OOF |
| Random seed | 42 |

## References


[1] N. Siddiqi, *Credit Risk Scorecards: Developing and Implementing Intelligent Credit Scoring*. Wiley, 2006.

[2] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

[3] I. Covert, S. M. Lundberg, and S.-I. Lee, "Understanding global feature contributions with additive importance measures (SAGE)," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

[4] B. D. Williamson and J. Feng, "Efficient nonparametric statistical inference on population feature importance using Shapley values," in *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020.

[5] M. B. Kursa and W. R. Rudnicki, "Feature selection with the Boruta package," *Journal of Statistical Software*, vol. 36, no. 11, pp. 1–13, 2010.

[6] N. Meinshausen and P. Bühlmann, "Stability selection," *Journal of the Royal Statistical Society: Series B*, vol. 72, no. 4, pp. 417–473, 2010.

[7] L. Breiman, J. H. Friedman, R. A. Olshen, and C. J. Stone, *Classification and Regression Trees*. Wadsworth, 1984.

[8] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, 2nd ed. Springer, 2009.

[9] E. R. DeLong, D. M. DeLong, and D. L. Clarke-Pearson, "Comparing the areas under two or more correlated receiver operating characteristic curves: a nonparametric approach," *Biometrics*, vol. 44, no. 3, pp. 837–845, 1988.

[10] S. Lessmann, B. Baesens, H.-V. Seow, and L. C. Thomas, "Benchmarking state-of-the-art classification algorithms for credit scoring: An update of research," *European Journal of Operational Research*, vol. 247, no. 1, pp. 124–136, 2015.

[11] Federal Reserve, OCC, and FDIC, "Supervisory guidance on model risk management (SR 26-2)," Interagency guidance, U.S. prudential regulators, 2026, supersedes SR 11-7 (2011).

[12] Basel Committee on Banking Supervision, "CRE: Calculation of RWA for credit risk," Bank for International Settlements, Basel Framework.

[13] European Parliament and Council, "Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (AI act)," 2024.

[14] M. Mitchell, S. Wu, A. Zaldivar, P. Barnes, L. Vasserman, B. Hutchinson, E. Spitzer, I. D. Raji, and T. Gebru, "Model cards for model reporting," in *Proc. Conf. on Fairness, Accountability, and Transparency (FAT*)*, 2019, pp. 220–229.

[15] T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, and K. Crawford, "Datasheets for datasets," *Communications of the ACM*, vol. 64, no. 12, pp. 86–92, 2021.

[16] U.S. Equal Employment Opportunity Commission, "Uniform guidelines on employee selection procedures," 29 C.F.R. Part 1607, 1978.

[17] M. Hardt, E. Price, and N. Srebro, "Equality of opportunity in supervised learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2016.

[18] P. Peduzzi, J. Concato, E. Kemper, T. R. Holford, and A. R. Feinstein, "A simulation study of the number of events per variable in logistic regression analysis," *Journal of Clinical Epidemiology*, vol. 49, no. 12, pp. 1373–1379, 1996.

[19] M. van Smeden, K. G. M. Moons, J. A. H. de Groot *et al.*, "Sample size for binary logistic prediction models: beyond events per variable criteria," *Statistical Methods in Medical Research*, vol. 28, no. 8, pp. 2455–2474, 2019.

[20] R. D. Riley, K. I. E. Snell, J. Ensor *et al.*, "Minimum sample size for developing a multivariable prediction model: Part ii — binary and time-to-event outcomes," *Statistics in Medicine*, vol. 38, no. 7, pp. 1276–1296, 2019.

[21] R. D. Riley, J. Ensor, K. I. E. Snell *et al.*, "Calculating the sample size

required for developing a clinical prediction model," *BMJ*, vol. 368, p. m441, 2020.

[22] R. D. Riley, B. Van Calster, and G. S. Collins, "A note on estimating the Cox-Snell $R^2$ from a reported C statistic (AUROC) to inform sample size calculations for developing a prediction model with a binary outcome," *Statistics in Medicine*, vol. 40, no. 4, pp. 859–864, 2021.

[23] D. J. Hand and W. E. Henley, "Statistical classification methods in consumer credit scoring: a review," *Journal of the Royal Statistical Society: Series A*, vol. 160, no. 3, pp. 523–541, 1997.

[24] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining (KDD)*, 2016, pp. 785–794.

[25] Y. Benjamini and Y. Hochberg, "Controlling the false discovery rate: a practical and powerful approach to multiple testing," *Journal of the Royal Statistical Society: Series B*, vol. 57, no. 1, pp. 289–300, 1995.

[26] J. Cohen, *Statistical Power Analysis for the Behavioral Sciences*, 2nd ed. Lawrence Erlbaum Associates, 1988.

[27] B. Zadrozny and C. Elkan, "Transforming classifier scores into accurate multiclass probability estimates," in *Proc. 8th ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining (KDD)*, 2002, pp. 694–699.

[28] A. Niculescu-Mizil and R. Caruana, "Predicting good probabilities with supervised learning," in *Proc. 22nd Int. Conf. on Machine Learning (ICML)*, 2005, pp. 625–632.

[29] H. Hofmann, "Statlog (German credit data)," UCI Machine Learning Repository, 1994.

[30] D. Dua and C. Graff, "UCI machine learning repository," University of California, Irvine, School of Information and Computer Sciences, 2019, http://archive.ics.uci.edu/ml.

[31] I.-C. Yeh and C.-h. Lien, "The comparisons of data mining techniques for the predictive accuracy of probability of default of credit card clients," *Expert Systems with Applications*, vol. 36, no. 2, pp. 2473–2480, 2009.

[32] B. Baesens, D. Rösch, and H. Scheule, *Credit Risk Analytics: Measurement Techniques, Applications, and Examples in SAS*. Wiley, 2016.

[33] Kaggle, "Give me some credit," Kaggle competition dataset, 2011, https://www.kaggle.com/c/GiveMeSomeCredit.

[34] FICO, "Explainable machine learning challenge (HELOC dataset)," FICO Community, 2018, https://community.fico.com/s/explainable-machine-learning-challenge.

[35] LendingClub, "LendingClub loan data," Public dataset, accessed via public mirrors.

[36] PAKDD, "PAKDD 2009 data mining competition: credit risk assessment," Pacific-Asia Conf. on Knowledge Discovery and Data Mining, 2009.

[37] M. P. Naeini, G. F. Cooper, and M. Hauskrecht, "Obtaining well calibrated probabilities using Bayesian binning," in *Proc. 29th AAAI Conf. on Artificial Intelligence*, 2015, pp. 2901–2907.

[38] D. W. Hosmer and S. Lemeshow, "Goodness of fit tests for the multiple logistic regression model," *Communications in Statistics — Theory and Methods*, vol. 9, no. 10, pp. 1043–1069, 1980.